

LDU: A Lightweight Concurrent Update Method with Deferred Processing for Linux Kernel Scalability

Joohyun Kyong and Sung-Soo Lim

School of Computer Science
Kookmin University


March 3, 2016

Outline

- ▶ Background of research and History of the Linux scalability
- ▶ Linux Scalability Problems
- ▶ Our new method and Evaluation
- ▶ Future plans and Summary

40 Years of Microprocessor Trend Data


40 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

40 Years of Microprocessor Trend Data

40 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

History of the Linux scalability research

Boyd-Wickizer, Silas, et al. "An Analysis of Linux Scalability to Many Cores." OSDI, 2010.

History of the Linux scalability research

Boyd-Wickizer, Silas, et al. "An Analysis of Linux Scalability to Many Cores." OSDI, 2010.

Austin T. Clements., M. Frans Kaashoek, and Nickolai Zeldovich.
"Scalable address spaces using RCU balanced trees." ASPLOS, 2012.


Austin T. Clements, M. Frans Kaashoek, Nickolai Zeldovich,
Robert T. Morris, Eddie Kohler.
"The Scalable Commutativity Rule: Designing Scalable Software for
Multicore Processors," SOSP 2013.

Austin T. Clements, M. Frans Kaashoek, and Nickolai Zeldovich.
"RadixVM: Scalable address spaces for multithreaded applications"
, EuroSys 2013

What about update?

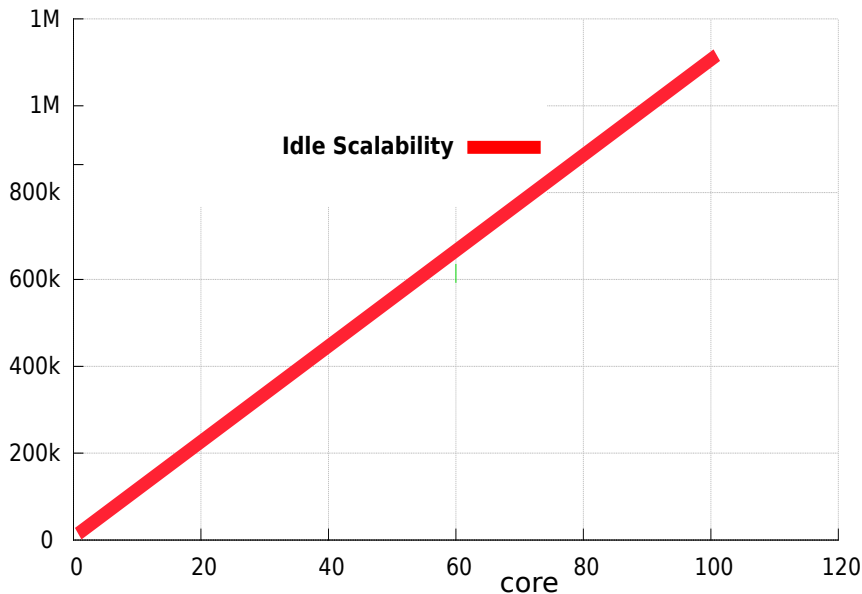
Linux Scalability Problems – AIM7– multiuser

jobs/min



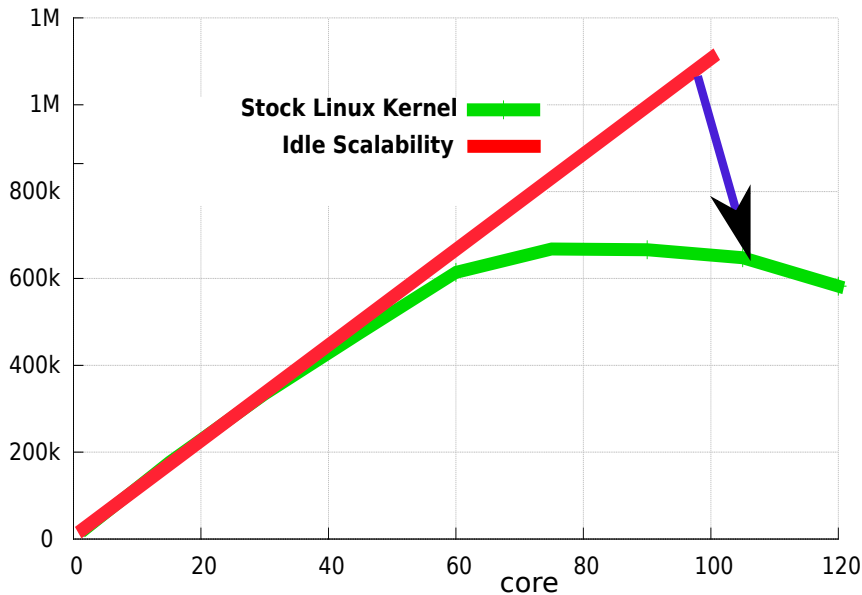
Linux Scalability Problems

jobs/min




Linux Scalability Problems

jobs/min




Linux Scalability Problems

jobs/min

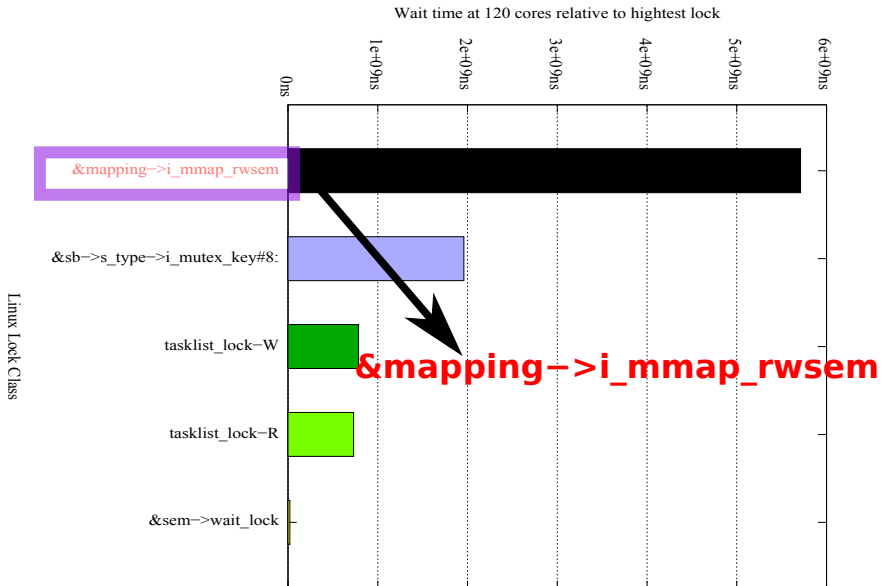


Linux Scalability Problems

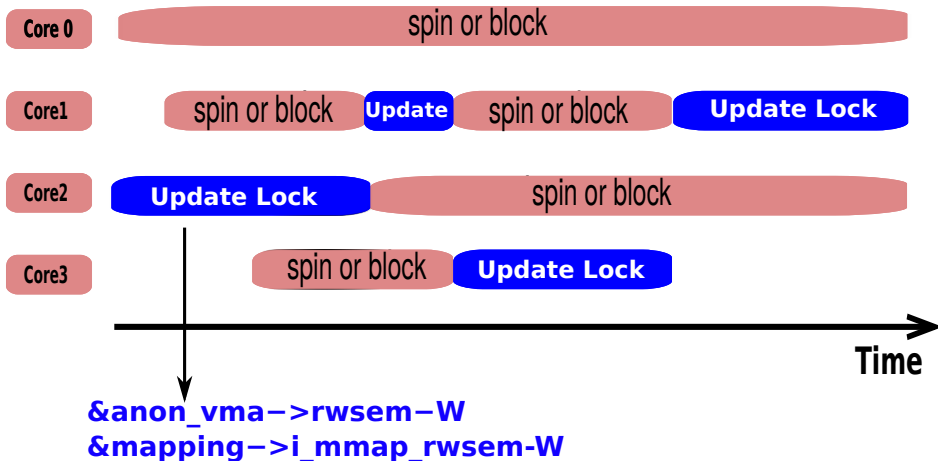
Linux Lock Class




Linux Scalability Problems




Update Serialization



Solution – Non blocking Data Structure





Solution – Non blocking Data Structure





- Harris' linked list : linked list
- Fraser skip list : skip list
- Lock-free tree : binary tree
- Michael's hash table : hash table
- Cliff Click's hash map : hash table
- No hot spot skip list : skip list
- Rotating skip list : skip list

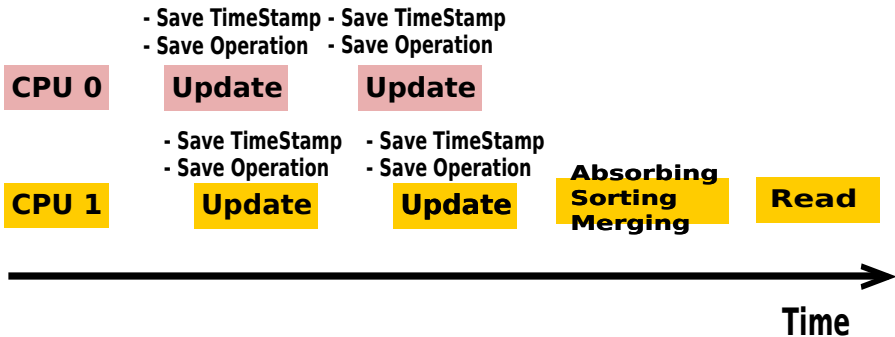

Per-core processing



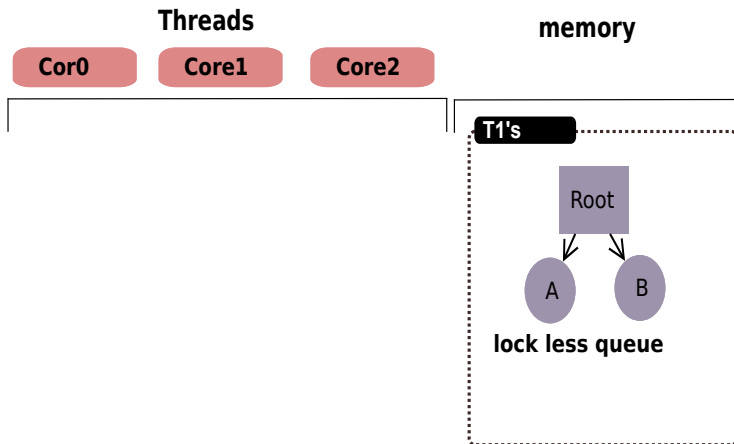
Per-core processing



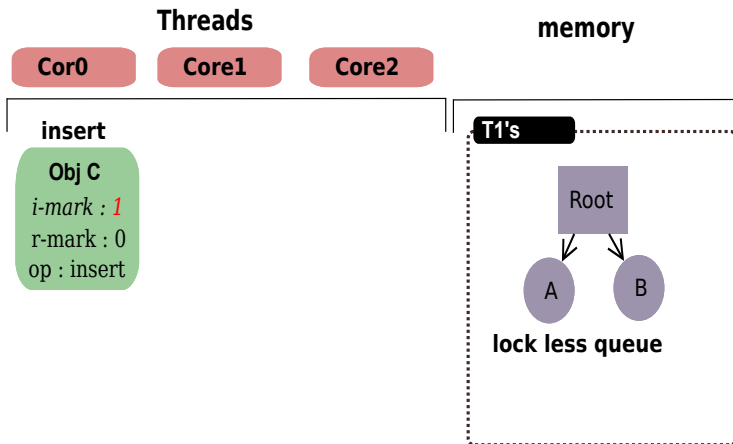
Per-core processing



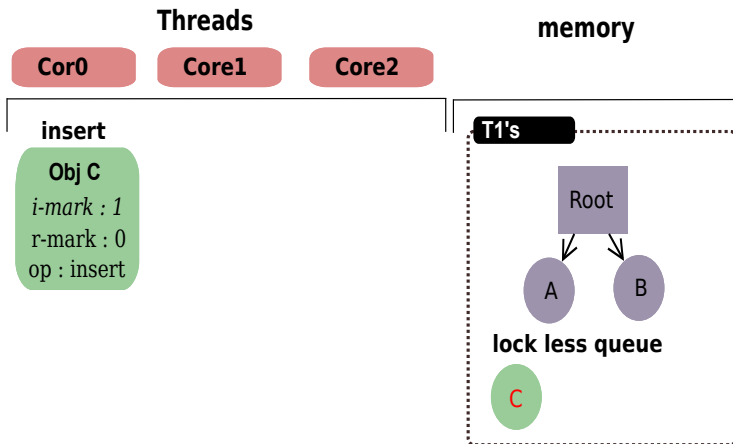
LDU: An example of Lightweight Deferred Update



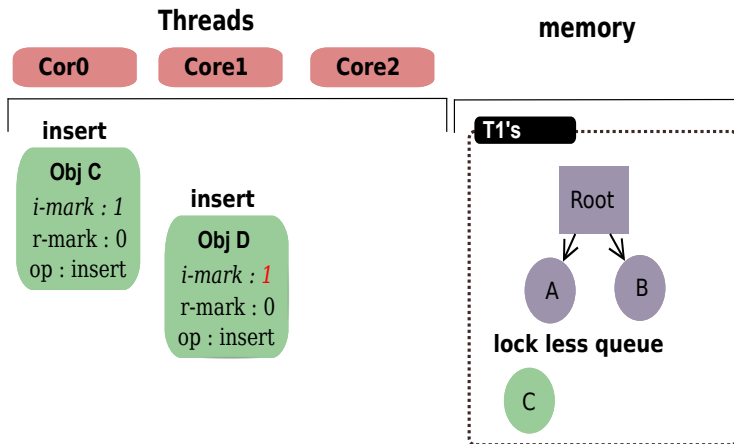
LDU: An example of Lightweight Deferred Update



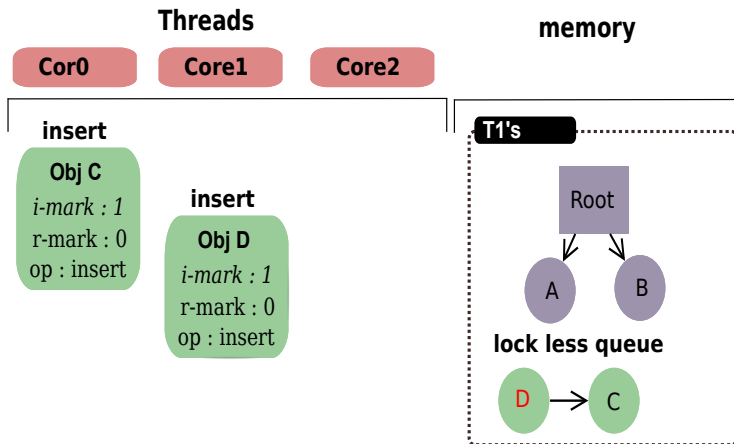
LDU: An example of Lightweight Deferred Update



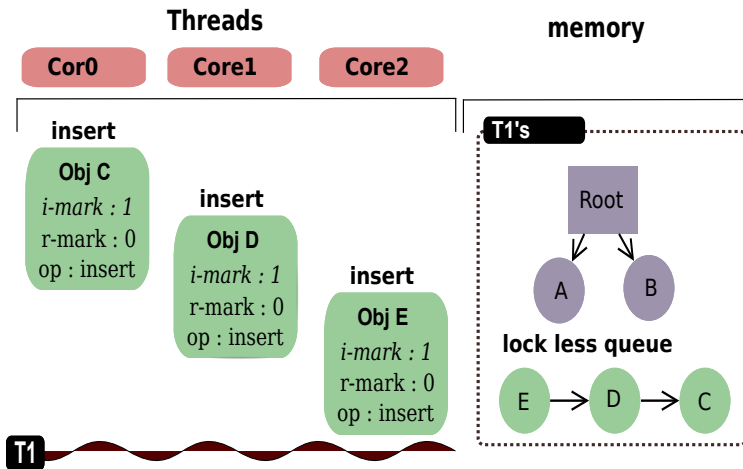
LDU: An example of Lightweight Deferred Update



LDU: An example of Lightweight Deferred Update



LDU: An example of Lightweight Deferred Update



LDU: An example of Lightweight Deferred Update

remove

Obj C

i-mark : 1

r-mark : 0

op : insert

T2's

Root



lock less queue



LDU: An example of Lightweight Deferred Update

remove

Obj C


i-mark : 0

r-mark : 0

op : insert

Update-side Absorbing

T2's



LDU: An example of Lightweight Deferred Update

remove

Obj C

i-mark : 0

r-mark : 0

op : insert

remove


Obj A

i-mark : 0

r-mark : 1

op : remove

T2's



LDU: An example of Lightweight Deferred Update

remove

Obj C

i-mark : 0
r-mark : 0
op : insert

remove

Obj E

i-mark : 0
r-mark : 0
op : insert

remove

Obj A

i-mark : 0
r-mark : 1
op : remove

T2

Tree lock

T2's

Root

A

B

lock less queue

A


E

D

C


LDU: An example of Lightweight Deferred Update

**Tree lock
synchronize**



LDU: An example of Lightweight Deferred Update

**Tree lock
synchronize**



LDU: An example of Lightweight Deferred Update

**Tree lock
synchronize**

physical
update 

physical
update 

T3

T3's

Root


D

B


lock less queue

NULL

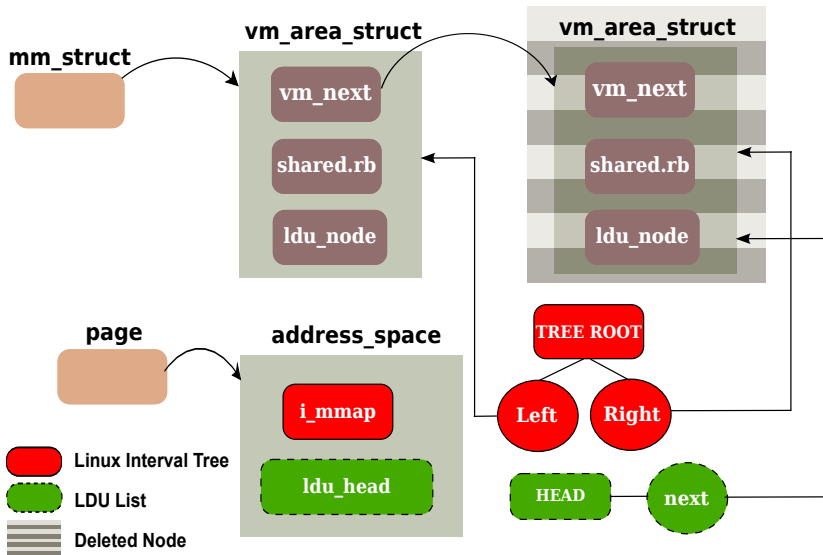
LDU: An example of Lightweight Deferred Update



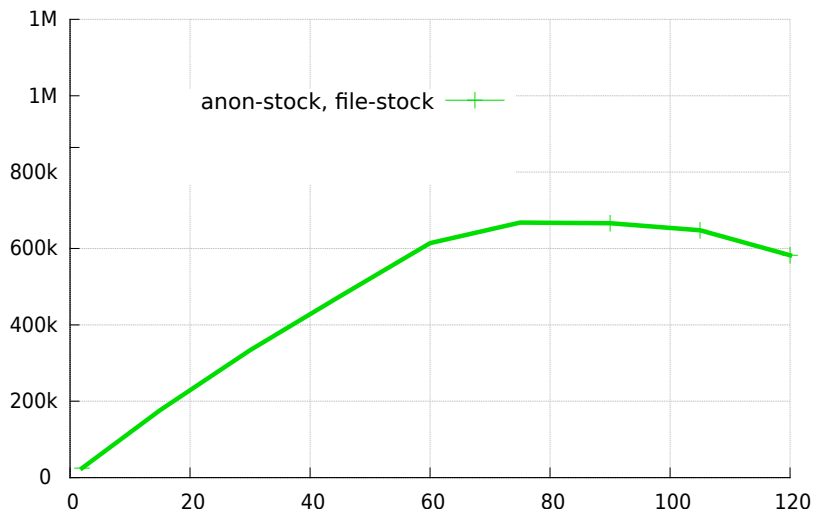
Concurrent update for Linux kernel




Concurrent update for Linux kernel




Evaluation – AIM7




Evaluation – AIM7




Evaluation – AIM7




Evaluation – AIM7




Evaluation – EXIM




Evaluation – EXIM



Evaluation – EXIM



Evaluation – EXIM



Summary

- ▶ History of the Linux scalability
- ▶ Linux Scalability Problems
- ▶ LDU and Evaluation
- ▶ Future plans and Summary
- ▶ <https://github.com/KMU-embedded/scalablelinux>

Q n A

Q & A